

Betsy -- Commit Timeline

Chronological build history mapped to project phases and decision logs.

Phase 0 -- Environment setup

2525b43 -- Initial commit: Betsy procurement agent test environment Starting point: mock datasets (12 SKUs, 6 suppliers, 15 invoices), FastAPI server, 4 injectable scenarios, and a live procurement dashboard. Everything the agent would need to test against -- before any agent code existed. -> DL-01

f684cf8 -- Add BPM analysis: before/after process model for Betsy agent Standalone HTML document (bpm_analysis.html) with swimlane diagrams showing the AS-IS manual procurement workflow vs the TO-BE agentic AI-augmented process. Includes 3- and 4-lane diagrams, scenario walkthroughs (stockout, duplicate invoice, price spike), and a weekly time-motion analysis. -> Pre-DL-02 / supports DL-01 framing

Phase 1 -- Agent architecture

c722d87 -- Add agent architecture: Pipeline, Orchestra, compare runner, and shared utilities Both agent architectures committed together: LangGraph Pipeline (6-node sequential) and Orchestra (3 parallel agents + coordinator). Shared utilities (llm.py, api_client.py, state.py) written once and used by both. -> DL-02

1daaba0 (merge) -- Merge pull request #1 from YashAlwani/agent-architecture Architecture branch merged into main.

9fd7ab7 -- Add analysis segments and decision logs (DL-01, DL-02) DL-01 and DL-02 written and committed alongside the architecture code. Evidence files (inventory.json, suppliers.json, graph.py, decide.py, PIPELINE_ARCHITECTURE.md, ORCHESTRA_ARCHITECTURE.md) added to decision_logs/evidence/. -> DL-01 · DL-02

Phase 2 -- UI and live integration

a737cac -- remove: delete Streamlit dashboard, replaced by betsy.html Streamlit prototype removed. The Flask/Jinja dashboard was replaced by a plain HTML approach served directly from FastAPI -- no build step, no framework dependency.

8d9f213 -- feat: add Betsy AI layer UI, docs, model switch to llama3.1:8b betsy.html wired to the live FastAPI endpoints. POST /api/run-agent added as a background task trigger. Model switched from Mistral 7B to Llama 3.1 8B for better structured JSON output. AI narrative, ? Risk column, and ? Betsy Score column all pulling from the live API. -> DL-04

4282f35 -- add DL-03, DL-04, DL-SUMMARY and all screenshots/evidence DL-03 (UI design) and DL-04 (implementation and first real run) written. All screenshots for both DLs added to decision_logs/images/. DL-SUMMARY first version committed. -> DL-03 · DL-04

Phase 3 -- HITL approval queue

2733e53 -- feat: DL-05 HITL approval queue -- full approve/decline loop /api/approvals router built (GET pending, POST approve, POST reject). act.py stores the full PO payload at queue time and executes it on approval. betsy.html approval cards wired to real endpoints. -> DL-05

e28613d -- fix: don't clear approvals on scenario reset -- approvals survive until user acts Bug found during DL-05 testing: state.reset() was wiping the approval queue at the end of every pipeline run. Single-line fix -- approvals kept independent of scenario state. -> DL-05

457f41a -- docs: DL-05 test report -- 11/11 passed, bug fix documented Test report for the 11-case approval queue test suite committed to docs/test-report-dl05.md. Bug and fix documented within the report. -> DL-05

f78a75f -- docs: write DL-05 -- HITL approval queue, bug doc, integration notes DL-05 decision log written. -> DL-05

ef39d1a (merge) -- merge: DL-05 HITL approval queue

Phase 4 -- Persistence, learning, and scheduling

fba4d06 -- feat: DL-06 persistence, EMA learning, APScheduler, stats panel Three gaps closed in one sprint: SQLite persistence via server/db.py (agent_log + approvals survive restarts), EMA supplier scoring triggered on PO delivery ($\tau=0.2$), APScheduler BackgroundScheduler firing the pipeline every 30 minutes, stats endpoint + betsy.html stats panel. -> DL-06

3df6380 -- docs: DL-06 test report -- 9/9 passed, EMA formula verified Test report committed. Restart verification (T4) and both EMA formula checks (T5a, T5b) documented with exact predicted vs actual values. -> DL-06

5ccea0 -- feat: add preserve_scores, score PATCH endpoint, magnus opus test, run_server script PATCH endpoint for supplier scores added to server/routers/suppliers.py. preserve_scores mechanism added for test isolation. Long-term learning integration test (test_long_term_learning.py) built -- two real LLM pipeline runs with 8 delivery simulation rounds between them. run_server.py script added for convenience. -> DL-07

9b8e7f7 -- docs: add DL-06 and DL-07 decision logs DL-06 and DL-07 written. -> DL-06 · DL-07

0d7dedb -- polish all DLs: remove inline code blocks, replace with casual file references Formatting pass across all DLs -- inline code blocks replaced with plain file references for readability.

cdaad08 -- docs: add README

bf8bd10 (merge) -- Merge pull request #2 from YashAlwani/feat/dl06-persistence-learning

Phase 5 -- Push notifications

6cdd95e -- feat: add desktop and email push notifications (DL-08) Two-channel notification system added from scratch -- no predecessor code existed. Desktop toasts via plyer (optional dependency), HTML emails via smtplib (stdlib). All logic in

server/notifier.py. Settings panel added to betsy.html. Runtime config via POST /api/notifications/config backed by mutable module-level attributes in server/config.py. 22/22 unit tests passing. -> DL-08

At a glance

Phase	Key deliverable	DL
Environment	Mock data + FastAPI server + dashboard	DL-01
Analysis	BPM as-is / to-be swimlane diagrams	--
Architecture	Pipeline + Orchestra + shared utilities	DL-02
UI	betsy.html live integration, model switch	DL-03 · DL-04
HITL	Approval queue end-to-end, bug fix, 11/11 tests	DL-05
Persistence	SQLite + EMA scoring + APScheduler	DL-06
Learning	Long-term learning integration test	DL-06 · DL-07
Notifications	Desktop + email push, settings panel, 22 tests	DL-08